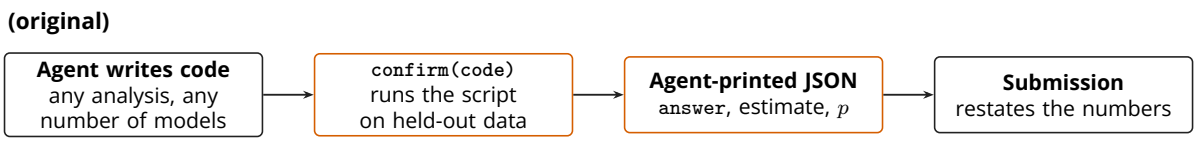
v1 (original)
Agent writes code
any analysis, any number of models
confirm(code)
runs the script on held-out data
Agent-printed JSON
answer, estimate, p
Submission
restates the numbers
v2 (trusted)
Agent names one
candidate exposure
confirm(exposure)
locks H; no code accepted
Harness OLS test
hash-bound evidence record
Submission
cites analysis_id