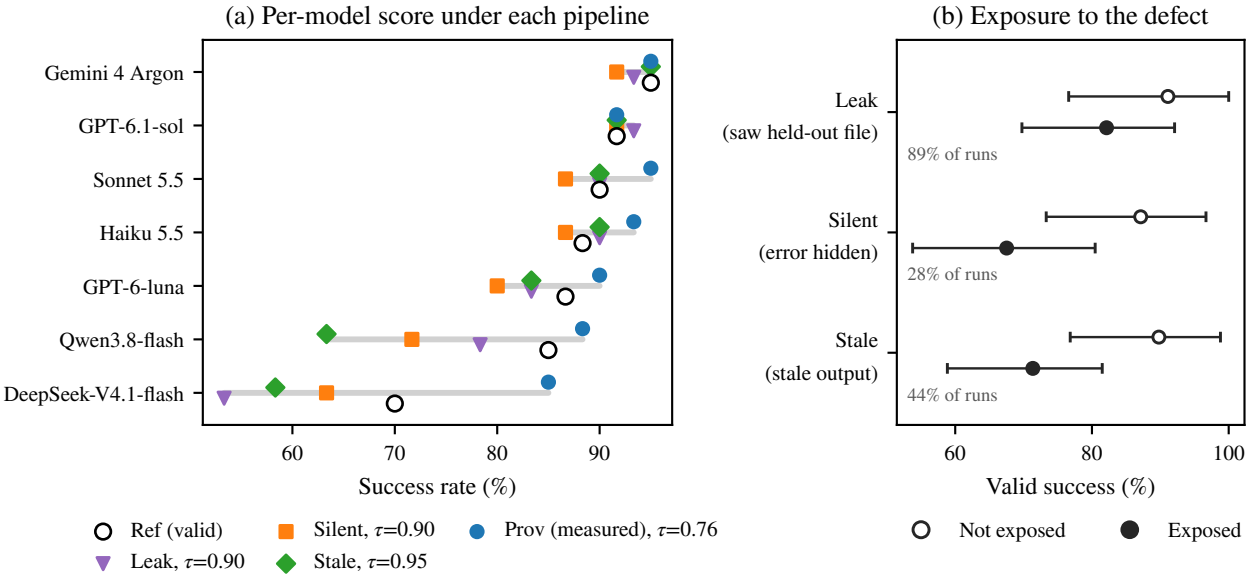
(a) Per-model score under each pipeline
Gemini 4 Argon
GPT-6.1-sol
Sonnet 5.5
Haiku 5.5
GPT-6-luna
Qwen3.8-flash
DeepSeek-V4.1-flash
60
70
80
90
Success rate (%)
Ref (valid)
Silent, τ=0.90
Prov (measured), τ=0.76
Leak, τ=0.90
Stale, τ=0.95
(b) Exposure to the defect
Leak
(saw held-out file)
89% of runs
Silent
(error hidden)
28% of runs
Stale
(stale output)
44% of runs
60
80
100
Valid success (%)
Not exposed
Exposed